

Lecture 3:

CNN: Back-propagation

boris.ginzburg@intel.com

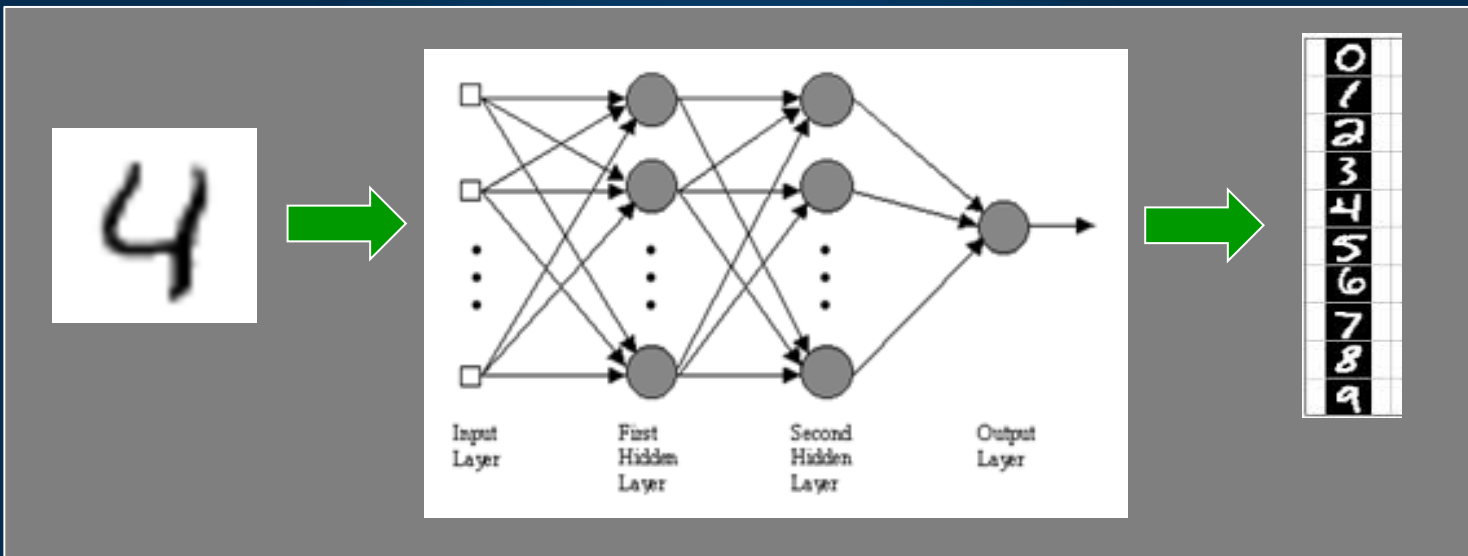
Agenda

- Introduction to gradient-based learning for Convolutional NN
- Backpropagation for basic layers
 - Softmax
 - Fully Connected layer
 - Pooling
 - ReLU
 - Convolutional layer
- Implementation of back-propagation for Convolutional layer
- CIFAR-10 training

Good Links

1. <http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>
2. <http://www.iro.umontreal.ca/~pift6266/H10/notes/gradient.html#flowgraph>

Gradient based training



Conv. NN is a function $y = f(x_0, w)$, where

x_0 is image [28,28],

w – network parameters (weights, bias)

y – softmax output= probability that x belongs to one of 10 classes 0..9

Gradient based training

We want to find parameters W , to minimize an error

$$E(f(x_0, w), y_0) = -\log(f(x_0, w) - y_0).$$

For this we will do iterative gradient descent:

$$w(t) = w(t-1) - \lambda * \frac{-\partial E}{\partial w}(t)$$

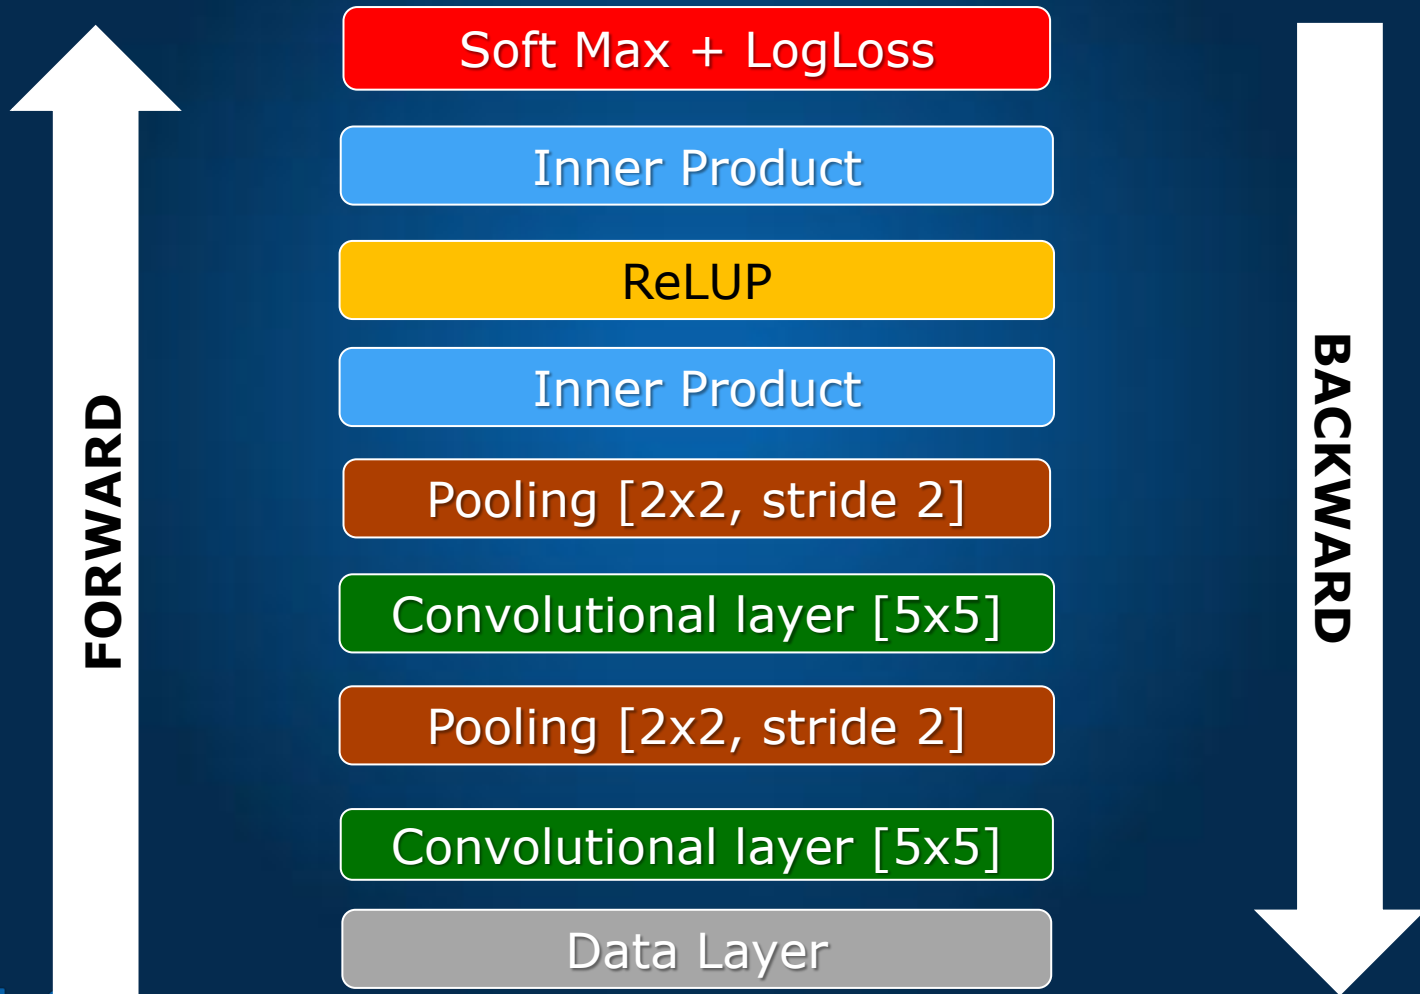
How do we compute gradient of E wrt weights?

Loss function E is chain of functions. Let's go layer by layer, from last layer back, and use the chain rule for gradient of complex functions:

$$\frac{\partial E}{\partial y_{l-1}} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial y_{l-1}}$$

$$\frac{\partial E}{\partial w_l} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial w_l}$$

LeNet topology



Layer:: Backward()

```
class Layer {  
    Setup (bottom, top);    // initialize layer  
    Forward (bottom, top);  //compute :  $y_l = f(w_l, y_{l-1})$   
    Backward( top, bottom); //compute gradient  
}
```

Backward: we start from gradient $\frac{\partial E}{\partial y_l}$ from last layer and

1) propagate gradient back : $\frac{\partial E}{\partial y_l} \rightarrow \frac{\partial E}{\partial y_{l-1}}$

2) compute the gradient of E wrt weights w_l : $\frac{\partial E}{\partial w_l}$

Softmax with LogLoss Layer

Consider the last layer, softmax with log-loss (MNIST example):

$$E = -\log(p_{y_0}) = -\log\left(\frac{e^{y_0}}{\sum_0^9 e^{y_k}}\right) = -y_0 + \log\left(\sum_0^9 e^{y_k}\right)$$

For all $k=0..9$, except k_0 (right answer) we want to decrease p_k :

$$\frac{\partial E}{\partial y_k} = \frac{e^{y_k}}{\sum_0^9 e^{y_k}} = p_k ,$$

for $k=k_0$ (right answer) we want to increase p_k :

$$\frac{\partial E}{\partial y_{k_0}} = -(1 - p_{k_0})$$

See http://ufldl.stanford.edu/wiki/index.php/Softmax_Regression

Inner product (Fully Connected) Layer

Fully connected layer is just Matrix - Vector multiplication:

$$y_l = W_l * y_{l-1}$$

So
$$\frac{\partial E}{\partial y_{l-1}} = \frac{\partial E}{\partial y_l} * W_l^T$$

and
$$\frac{\partial E}{\partial W_l} = \frac{\partial E}{\partial y_l} * y_{l-1}$$

Note: we need y_{l-1} , so we should keep them on forward pass.

ReLU Layer

Rectified Linear Unit :

$$y_l = \max (0, y_{l-1})$$

$$\text{SO} \quad \frac{\partial L}{\partial y_{l-1}} = \begin{cases} = 0, & \text{if } (y_l < 0) \\ = \frac{\partial L}{\partial y_l}, & \text{otherwise} \end{cases}$$

Max-Pooling Layer

Forward :

for (p = 0; p < k; p++)

for (q = 0; q < k; q++)

$y_n(x, y) = \max(y_n(x, y), y_{n-1}(x + p, y + q));$



Backward:

$$\frac{\partial L}{\partial y_{n-1}}(x + p, y + q) = \begin{cases} = 0, & \text{if } (y_n(x, y) \neq y_{n-1}(x + p, y + q)) \\ = \frac{\partial L}{\partial y_n}(x, y), & \text{otherwise} \end{cases}$$


Quiz:

1. What will be gradient for Sum-pooling?
2. What will be gradient if pooling areas overlap? (e.g. stride = 1)?

Convolutional Layer :: Backward

```

for (n = 0; n < N; n++)
  for (m = 0; m < M; m++)
    for (y = 0; y < Y; y++)
      for (x = 0; x < X; x++)
        for (p = 0; p < K; p++)
          for (q = 0; q < K; q++)
            yL(n; x, y) += yL-1(m, x+p, y+q) * w(n, m; p, q);
    
```



Let' s use the chain rule for convolutional layer


$$\frac{\partial E}{\partial y_{l-1}} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial y_{l-1}};$$

$$\frac{\partial E}{\partial w_l} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial w_{l-1}}$$

Convolutional Layer :: Backward

Example: $M=1$, $N=2$, $K=2$.

Take one pixel in level $(n-1)$. Which pixels in next level are influenced by it?



Convolutional Layer :: Backward

Let's use the chain rule for convolutional layer:

Gradient $\frac{\partial E}{\partial y_{l-1}}$ is sum of convolution with gradients $\frac{\partial E}{\partial y_l}$ over all feature maps from "upper" layer:

$$\frac{\partial E}{\partial y_{l-1}} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial y_{l-1}} = \sum_{n=1}^N \text{back_corr}(W, \frac{\partial E}{\partial y_l})$$

Gradient of E wrt w is sum over all "pixels" (x,y) in the input map :

$$\frac{\partial E}{\partial w_l} = \frac{\partial E}{\partial l} \times \frac{\partial y_l(w, y_{l-1})}{\partial w_l} = \sum_{\substack{0 \leq x \leq X \\ 0 < y < Y}} \left(\frac{\partial E}{\partial y_l}(x, y) \circ y_{l-1}(x, y) \right)$$

Convolutional Layer :: Backward

How this is implemented:

```
backward ( ) { ...
```

```
// im2col data to col_data
```

```
im2col_cpu( bottom_data , CHANNELS_ HEIGHT_ WIDTH_ KSIZE_ PAD_ STRIDE_  
            col_data);
```

```
// gradient w.r.t. weight.:
```

```
caffe_cpu_gemm (CblasNoTrans, CblasTrans, M_ K_ N_ 1., top_diff, col_data , 1.,  
               weight_diff );
```

```
// gradient w.r.t. bottom data:
```

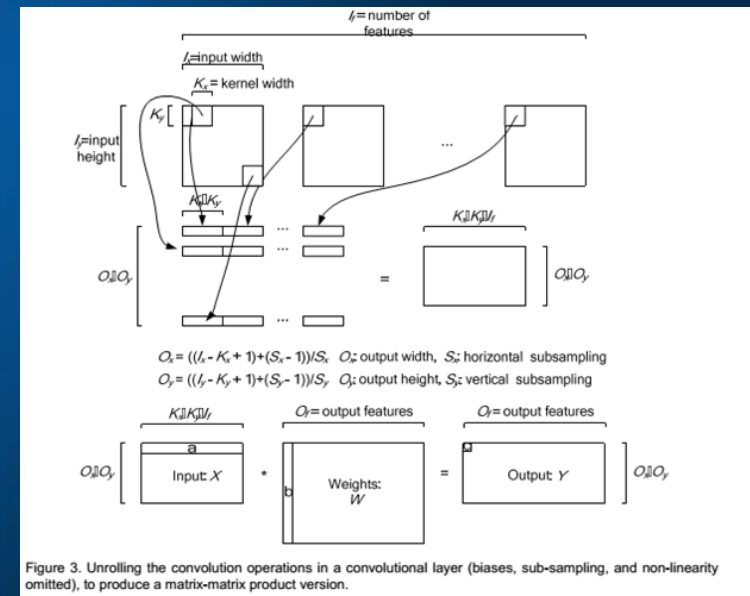

```
caffe_cpu_gemm (CblasTrans, CblasNoTrans, K_ N_ M_ 1., weight , top_diff , 0.,  
               col_diff );
```

```
// col2im back to the data
```

```
col2im_cpu(col_diff, CHANNELS_ HEIGHT_ WIDTH_ KSIZE_ PAD_ STRIDE_  
           bottom_diff );
```

Convolutional Layer : im2col

Implementation is based on reduction of convolution layer to matrix – matrix multiply (See Chellapilla et al , “High Performance Convolutional Neural Networks for Document Processing”)




CIFAR-10 Training

<http://www.cs.toronto.edu/~kriz/cifar.html>

<https://www.kaggle.com/c/cifar-10>


60000 32x32 colour images in 10 classes, with 6000 images per class.
There are:

- 50000 training images
- 10000 test images.



Exercises

1. Look at definition of following layers (Backward)
 - sigmoid, tanh
2. Implement a new layer:
 - softplus $y_l = \log(1 + e^{y_{l-1}})$
3. Train CIFAR-10 with different topologies



Project:

1. Port CIFAR-100 to caffe